

Vector Processing

Muhammad Tahir

Lecture 17

Electrical Engineering Department
University of Engineering and Technology Lahore

Contents

- 1 Parallelism in Execution
- 2 Data-Level Parallelism
- 3 RISC V Vector Extension
- 4 RV Vector Operations

Types of Parallelism

- **Instruction-Level Parallelism (ILP):** Execute multiple independent instructions from one stream (task, thread) in parallel (**pipelining, superscalar, VLIW**)
- **Data-Level Parallelism (DLP):** Execute multiple operations of the same type in parallel (**SIMD (vector, array)** processing)
- **Thread-Level Parallelism (TLP):** Execute multiple instructions from independent streams in parallel (**multithreading, multi-core**)

Flynn's Classification

Flynn's Classification		Data Stream	
		Single	Multiple
Instruction Stream	Single	SISD	SIMD (Vector)
	Multiple	MISD	MIMD (Multi-core)

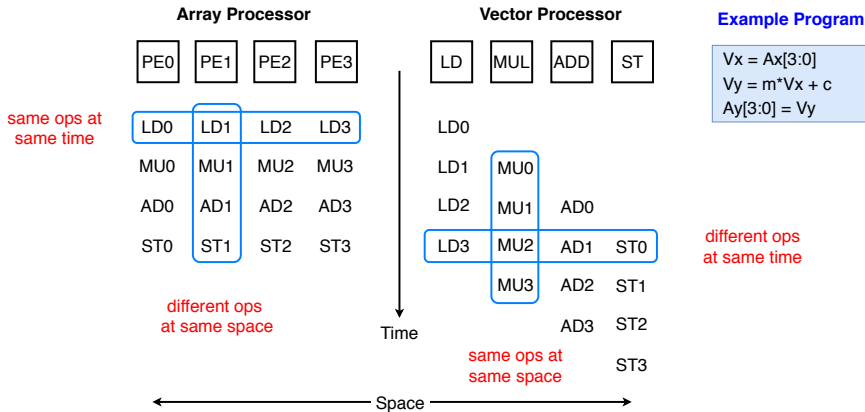
Advantages of SIMD

- Single 32-bit Microprocessor
 - One 32-bit element processing per instruction
- 512-bit SIMD Processor
 - 16 32-bit elements processing per instruction
 - 32 16-bit elements processing per instruction
 - 64 8-bit elements processing per instruction
- For 1 GHz SIMD processor clock frequency the performance is 64 GOPS (for 8-bit elements)

SIMD Processing

- *Single Instruction* operates on *Multiple Data* elements
- Multiple processing elements of same type
- Same instruction operates on multiple data elements at the same time using different spaces (**array processing**)
- Same instruction operates on multiple data elements in consecutive time steps using the same space (**vector processing**)
- *Array-Vector* duality is *Time-Space* duality

SIMD Processing: Vector vs Array





Example Program

```
Vx = Ax[3:0]
Vy = m * Vx + c
Ay[3:0] = Vy
```


Figure credit: Prof. Onur Mutlu

Vector Arithmetic




Vector Execution

Execution using
one pipelined
functional unit

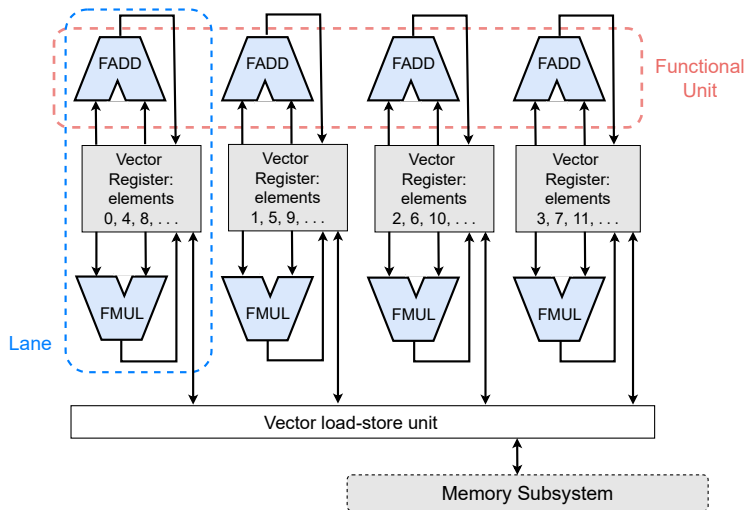


Time

Execution using
four pipelined
functional units



Structure of Vector Unit

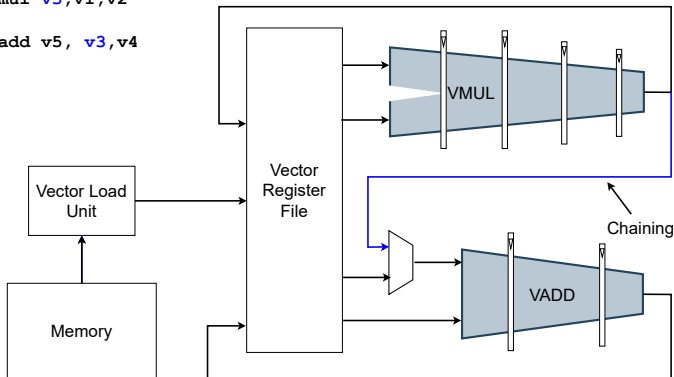


Vector Chaining

- Similar to register bypassing

```
vmul v3,v1,v2
```

```
vadd v5, v3,v4
```



Vector Stripmining

- Consider an application with data length, N
- Vector registers are of finite length VL_{max} and let $N > VL_{max}$
- **Stripmining**: Break the loop into segments that fit in vector registers
- For first iteration set

$$VL = N \bmod VL_{max}$$


- For second and onward iterations set

$$VL = VL_{max}$$


Vector Masking

- Require conditional execution of vector operations
- Introduce **vector mask register**
 - Tells which elements of a vector to be operated on
 - 1-bit masking per element
 - Maskable vector instructions
 - Vector operation becomes *bubble*, when mask bit is clear
- Vector masking implementation
 - **Density Time Implementation:** Scan mask vector and execute nonzero mask elements
 - **Simplified Implementation:** Perform operation on all elements, but control each result writeback based on mask bit

Vector Masking Cont'd



**Simplified
Implementaton**



**Density Time
Implementaton**

RISC V Vector Extension

- RISC V Scalar State
 - Program counter (PC) register
 - 32 XLEN-bit integer registers (x0-x31)
 - 32 FLEN-bit floating point (**optional**) registers (f0-f31)
 - Control and Status Registers (CSRs) (Privileged Architecture Specs.)
 - Floating point (**optional**) status register (fCSR) for rounding mode & exception status

RISC V Vector Extension Cont'd

- To support vector extension first define the following three parameters
 - Element Length (**ELEN**): Defines the maximum size of a single vector element in bits, $ELEN \geq 8$ and must be a power of 2
 - Vector Length (**VLEN**): Defines the number of bits in a vector register, $VLEN \geq ELEN$ and must be a power of 2
 - Stripping Length (**SLEN**): Defines the striping distance in bits, which must comply $VLEN \geq SLEN \geq 32$ and be a power of 2

RISC V Vector Registers

- Basic Requirements of vector register
 - Each vector data register holds N M-bit values
 - SEW – single vector element in bits, M-bit, power of 2
 - VLEN – number of bits in a vector register, N*M-bit, power of 2, $VLEN \geq SEW$
 - VLMAX – maximum number of vector elements, which can be operated on by single vector instruction (and is equal to $LMUL(VLEN/SEW)$)
 - There are 32 VLEN-bit vector data registers (v0-v31)
 - Vector control registers: VMASK (mask – predication – valid element for execution) is v0

RISC V Vector Registers Cont'd

- RISC V Vector Extension (state augmentation)
 - 32 VLEN-bit vector registers (v0-v31)
 - Vector length register **vl** to specify number of elements to be updated (of destination vector register) by a vector instruction
 - Vector type register **vtype**
 - Vector start register **vstart** for trap handling and specifies the index of first element to be operated by vector instruction
 - Vector byte length register **vlenb** (vector length in bytes)
 - Vector fixed-point rounding mode register **vxrm**
 - Vector fixed-point saturation flag register **vxsat**
 - Vector control and status register **vcsr** (combines **vxrm** and **vxsat** as bit-fields of this register)


Vector Extension CSRs

Table 1: CSRs for vector extension.

Address	Priv.	Reg. Name	Description
0x008	URW	vstart	Vector start position
0x009	URW	vxsat	Fixed-point saturate flag
0x00A	URW	vxrm	Fixed-point rounding mode
0x00F	URW	vcsr	Vector control and status register
0xC20	URO	vl	Vector length
0xC21	URO	vtype	Vector data type register
0xC22	URO	vlenb	VLEN/8 (vector register length in bytes)

vtype Register

- **vtype** (vector type) register bit field descriptions



Bit field	Description
vlmul[2:0]	Vector register group <i>length multiplier</i> (LMUL) setting
vsew[2:0]	Vector <i>standard element width</i> (SEW) setting
vta	Vector tail agnostic
vma	Vector mask agnostic
vill	Illegal value if set

vtype Register Cont'd

- The value in **vsew** sets the dynamic *Standard Element Width* (SEW)
- A vector register is viewed as being divided into $\frac{VLEN}{SEW}$ standard-width elements


vsew[2:0]	SEW = 8×2^{vsew}	Elements/vector register (VLEN = 1024 bits)
000	8	128
001	16	64
010	32	32
011	64	16
100	128	8
101	256	4
110	512	2
111	1024	1

vtype Register Cont'd

- The **vlmul** is a signed bit-field
- Length Multiplier* $LMUL = 2^{v_{lmul}}$ (for vector register group)
 - Allows multiple vector registers to be grouped together, with single vector instruction operating on multiple vector registers
 - Can be a fractional value, reducing the number of bits used in a vector register


vlmul[2:0]	LMUL = $2^{v_{lmul}}$	No. of groups	Indexed register grouping
100	-	-	Reserved
101 (= -3)	1/8	32	only v_n in the group
110 (= -2)	1/4	32	only v_n in the group
111 (= -1)	1/2	32	only v_n in the group
000	1	32	only v_n in the group
001	2	16	v_n, v_{n+1} in the group
010	4	8	v_n, \dots, v_{n+3} in the group
011	8	4	v_n, \dots, v_{n+7} in the group

vtype Register Cont'd



- LMUL = 2, instructions with odd register are illegal
- LMUL = 4, vector registers are incremented by 4, else illegal instruction
- LMUL = 8, only v0, v8, v16, and v24 are valid vector registers

vtype Register Cont'd




- LMUL=1/2, load data into half vector register, then perform vector double operations to expand data to the whole register
- LMUL=1/4, using vector extension of 4x to expand data to the whole register
- LMUL=1/8, using vector extension of 8x to expand data to the whole register


vtype Register Cont'd

- **vta** (*vector tail agnostic*) bit controls updating of destination tail elements by a vector instruction
- **vma** (*vector mask agnostic*) bit allows masking of destination elements by a vector instruction
- The **vill** bit tells that a vector instruction attempted to write an unsupported value to **vtype**
- If **vill** bit is set, then any attempt to execute a vector instruction that depends upon **vtype** will raise an illegal-instruction exception
- When the **vill** bit is set, the other XLEN-1 bits in **vtype** should be zero

Vector Length Register vl



- vl is the number of elements for vector operations
- vl is set to VLMAX if the set value is greater than VLMAX
- The default vl is VLMAX
- $vl = 0$, then no elements are updated in vector operations



VMASK Register

- v0 is used for vector mask, single mask for all LMUL values
- 1 mask bit for each element
 - 0 – mask (mask-off), no write to destination vector elements
 - 1 – unmask
- The mask is enabled for each vector instruction.
 - `vop.v v8, v16, 3, v0.t //` mask is enabled, each element writes back with `v0.mask[i] = 1`
 - `vop.v v8, v16, 3 //` unmask, all elements write result data back to VRF

RV V extension ISA

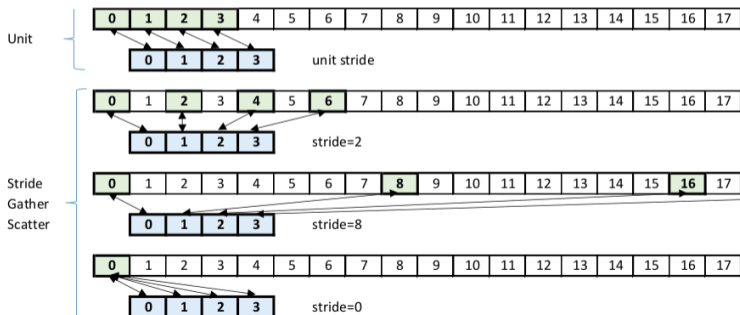
- RV V extension ISA includes:
 - Memory Operations
 - Compute Instructions
 - Permute Instructions

RVV ISA Memory Operations

- Load/store operations move groups of data between the VRF and memory
- Four basic types of addressing is used in RVV
 - Unit stride: fetch consecutive elements in contiguous memory
 - Non-unit or constant stride: elements are in stride memory; stride is the distance between 2 elements of a vector in memory
 - Indexed (gather-scatter): elements are randomly offset from a base address
 - Segment: move multiple contiguous fields in memory to and from consecutively numbered vector registers
- Note that Atomic load/store is not a part of vector ISA


RVV ISA Memory Operations

- Unit stride continuous memory access
- Constant stride: gather scatter (stride can be negative or zero)



RVV ISA Memory Operations

- The indices are absolute and not accumulative
- The indices are not in-order



Vector Load/Store Examples

- SEW=16, LMUL=1, VLEN=512b
- elements=512/16=32, all vector instructions operate on 32 elements
 - vle16.v v1, (t0) // Load a vector of sew = 16, lmul = 1, load to v1
 - vle32.v v2, (t1) // Load a vector of sew = 32, lmul = 2, load to v2 and v3
 - vse32.v v4, (t1) // Store a vector of sew = 32, emul = 2, v4 and v5 store to memory
 - vle8.v v1, (t1) // Load a vector of sew=8, emul = 1/2, load to elements first half of v1
 - vle64.v v4, (t0) // Load a vector of sew=64, emul = 4, load to v4, v5, v6, v7

RVV Compute Instructions

- RVV compute instructions includes:
 - Vector Arithmetic Instructions
 - Vector Widening Instructions
 - Vector Narrow Instructions
 - Vector-integer Arithmetic Instructions
 - Vector Bitwise Logical Instructions
 - Vector Bit Shift Instructions
 - Vector Integer Comparison Instructions
 - Vector Integer Min/Max Instructions
 - Vector Mask-Register Logical Instructions

Suggested Reading

- RISC-V Vector Extension Webinar II
[[RISC-V Vector Extension Webinar II](#)]
- Lecture slides on Vector extension [[Asanovic, 2020a](#)]
- Lecture slides on RISC V Vector extension [[Asanovic, 2020b](#)]
- RISC-V "V" Vector Extension Specifications, Draft Version 0.9 [[Asanovic et al., 2016](#)]
- RISC-V Vector Extension Webinar III
[[RISC-V Vector Extension Webinar III](#)]

References



Asanovic, K. (2020a).

“risc-v vectors, cs152, spring 2020”.

<http://inst.eecs.berkeley.edu/~cs152/sp20/lectures/L15-Vectors.pdf>.



Asanovic, K. (2020b).

“risc-v vectors, cs152, spring 2020”.

<http://inst.eecs.berkeley.edu/~cs152/sp20/lectures/L17-RISCV-Vectors.pdf>.



Asanovic, K., Waterman, A., and et. al. (2016).

The risc-v “v” vector extension specifications, draft version 0.9.



RISC-V Vector Extension Webinar II.

“RISC-V Vector Extension Webinar II, Andes Technology”.

http://www.andestech.com/wp-content/uploads/Andes-RVV-Webinar-II_final.pdf.



RISC-V Vector Extension Webinar III.

“RISC-V Vector Extension Webinar III, Andes Technology”.

http://www.andestech.com/wp-content/uploads/Andes-RVV-Webinar-III_final.pdf.